

Forecasting Unit Sales (Task 1)

```
In [98]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import xgboost as xgb
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split, GridSearchCV
import warnings
warnings.filterwarnings('ignore')

color_pal = sns.color_palette()
%matplotlib inline
```

```
In [99]: train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')
```

Step 1: Data Loading and Initial Inspection(Training Data)

```
In [100... train_df.head()
```

Out [100...

	ID	date	Item Id	Item Name	ad_spend	anarix_id	units
0	2022-04-12_B09KDTS4DC	2022-04-12	B09KDTS4DC	NapQueen Elizabeth 8" Gel Memory Foam Mattress...	NaN	NAPQUEEN	0.0
1	2022-04-12_B09MR2MLZH	2022-04-12	B09MR2MLZH	NapQueen 12 Inch Bamboo Charcoal Queen Size Me...	NaN	NAPQUEEN	0.0
2	2022-04-12_B09KSYL73R	2022-04-12	B09KSYL73R	NapQueen Elsa 8" Innerspring Mattress, Twin XL	NaN	NAPQUEEN	0.0
3	2022-04-12_B09KT5HMNY	2022-04-12	B09KT5HMNY	NapQueen Elsa 6" Innerspring Mattress, Twin	NaN	NAPQUEEN	0.0
4	2022-04-12_B09KTF8ZDQ	2022-04-12	B09KTF8ZDQ	NapQueen Elsa 6" Innerspring Mattress, Twin XL	NaN	NAPQUEEN	0.0

In [101...

```
train_df.tail()
```

Out [101...

	ID	date	Item Id	Item Name	ad_spend	anarix_id	uni
101485	2024-05-31_B0CR4BGLK5	2024-05-31	B0CR4BGLK5	NaN	604.73	NAPQUEEN	NaN
101486	2024-05-31_B0CR4BG4ZW	2024-05-31	B0CR4BG4ZW	NaN	261.21	NAPQUEEN	2
101487	2024-05-31_B0CR49NR3B	2024-05-31	B0CR49NR3B	NaN	0.00	NAPQUEEN	NaN
101488	2024-05-31_B0CR49N6MQ	2024-05-31	B0CR49N6MQ	NaN	0.00	NAPQUEEN	NaN
101489	2024-05-31_B0CR4BK4FW	2024-05-31	B0CR4BK4FW	NaN	0.00	NAPQUEEN	NaN

In [102...

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101490 entries, 0 to 101489
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              101490 non-null object
1   date            101490 non-null object
2   Item Id         101488 non-null object
3   Item Name       99658 non-null  object
4   ad_spend        77303 non-null  float64
5   anarix_id       101490 non-null object
6   units           83592 non-null  float64
7   unit_price      101490 non-null float64
dtypes: float64(3), object(5)
memory usage: 6.2+ MB
```

```
In [103... train_df.describe()
```

```
Out[103...
      ad_spend      units      unit_price
count  77303.000000  83592.000000  101490.000000
mean     110.771470    10.284381    106.750922
std      529.303777    68.945915    425.704733
min        0.000000   -173.000000   -8232.000000
25%        0.000000    0.000000    0.000000
50%        4.230000    1.000000    0.000000
75%       44.310000    5.000000    0.000000
max    47934.990000   9004.000000   21557.390000
```

```
In [104... train_df.isnull().sum()
```

```
Out[104...
ID              0
date            0
Item Id         2
Item Name      1832
ad_spend       24187
anarix_id       0
units          17898
unit_price      0
dtype: int64
```

```
In [105... print("Unique values in 'ID':", train_df['ID'].nunique())
print("Unique values in 'Item Id':", train_df['Item Id'].nunique())
print("Unique values in 'Item Name':", train_df['Item Name'].nunique())
```

```
Unique values in 'ID': 101490
Unique values in 'Item Id': 217
Unique values in 'Item Name': 199
```

Step 2: Data Loading and Initial Inspection(Testing Data)

In [106...

test_df.head()

Out [106...

	ID	date	Item Id	Item Name	ad_spend	anarix_id	unit_
0	2024-07-01_B09KDR64LT	2024-07-01	B09KDR64LT	NapQueen Elizabeth 10" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
1	2024-07-01_B09KDTS4DC	2024-07-01	B09KDTS4DC	NapQueen Elizabeth 8" Gel Memory Foam Mattress...	NaN	NAPQUEEN	
2	2024-07-01_B09KDTHJ6V	2024-07-01	B09KDTHJ6V	NapQueen Elizabeth 12" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
3	2024-07-01_B09KDQ2BWY	2024-07-01	B09KDQ2BWY	NapQueen Elizabeth 12" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
4	2024-07-01_B09KDYY3SB	2024-07-01	B09KDYY3SB	NapQueen Elizabeth 10" Gel Memory Foam Mattres...	101.72	NAPQUEEN	10

In [107...

test_df.tail()

Out [107...

	ID	date	Item Id	Item Name	ad_spend	anarix_id	u
2828	2024-07-28_B0BRCW2B64	2024-07-28	B0BRCW2B64	NapQueen Anula Green Tea 12", Queen	11.78	NAPQUEEN	
2829	2024-07-28_B0CFV6V981	2024-07-28	B0CFV6V981	NaN	1.17	NAPQUEEN	
2830	2024-07-28_B0BNL5BKMK	2024-07-28	B0BNL5BKMK	NapQueen 2" Bamboo Charcoal Mattress Topper, ...	0.00	NAPQUEEN	
2831	2024-07-28_B0CR49BQRS	2024-07-28	B0CR49BQRS	NaN	1.87	NAPQUEEN	
2832	2024-07-28_B0CY5QQ49F	2024-07-28	B0CY5QQ49F	NaN	1.45	NAPQUEEN	

In [108...

```
test_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2833 entries, 0 to 2832
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0    ID              2833 non-null   object
1    date            2833 non-null   object
2    Item Id         2833 non-null   object
3    Item Name       2489 non-null   object
4    ad_spend        1382 non-null   float64
5    anarix_id       2833 non-null   object
6    unit_price      2833 non-null   float64
dtypes: float64(2), object(5)
memory usage: 155.1+ KB
```

In [109...

```
test_df.describe()
```

	ad_spend	unit_price
count	1382.000000	2833.000000
mean	198.838032	98.725873
std	797.354508	383.585307
min	0.000000	-1988.180000
25%	0.730000	0.000000
50%	39.200000	0.000000
75%	156.012500	0.000000
max	18724.850000	6870.000000

```
In [110...] test_df.isnull().sum()
```

```
Out[110...] ID          0
date          0
Item Id       0
Item Name     344
ad_spend      1451
anarix_id     0
unit_price    0
dtype: int64
```

```
In [111...] print("\nUnique values in 'Item Id':", test_df['Item Id'].nunique())
print("Unique values in 'Item Name':", test_df['Item Name'].nunique())
```

```
Unique values in 'Item Id': 155
Unique values in 'Item Name': 142
```

STEP 3: DATA PRE PROCESSING

```
In [112...] #CLEAN TRAIN DATA
```

```
train_df['date'] = pd.to_datetime(train_df['date'])
train_df = train_df.sort_values(['Item Id', 'date'])
```

```
In [113...] train_df = train_df.dropna(subset=['Item Id'])
```

```
In [114...] train_df['ad_spend'] = train_df['ad_spend'].fillna(0)
train_df['units'] = train_df['units'].fillna(0)
```

```
In [115...] train_df['Item Name'] = train_df['Item Name'].fillna('Unknown')
```

```
In [116...] train_df['units'] = train_df['units'].clip(lower=0)
train_df['unit_price'] = train_df['unit_price'].clip(lower=0)
```

```
train_df['missing_name'] = train_df['Item Name'].apply(lambda x: 1 if x ==
```

```
In [117... print("Rows in dataset after cleaning:", len(train_df))
print("\nMissing values after cleaning:")
print(train_df.isnull().sum())
print("\nDataset info after cleaning:")
train_df.info()
```

Rows in dataset after cleaning: 101488

Missing values after cleaning:

```
ID          0
date         0
Item Id      0
Item Name    0
ad_spend     0
anarix_id    0
units        0
unit_price   0
missing_name 0
dtype: int64
```

Dataset info after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 101488 entries, 62967 to 101466
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	101488 non-null	object
1	date	101488 non-null	datetime64[ns]
2	Item Id	101488 non-null	object
3	Item Name	101488 non-null	object
4	ad_spend	101488 non-null	float64
5	anarix_id	101488 non-null	object
6	units	101488 non-null	float64
7	unit_price	101488 non-null	float64
8	missing_name	101488 non-null	int64

```
dtypes: datetime64[ns](1), float64(3), int64(1), object(4)
```

```
memory usage: 7.7+ MB
```

```
In [118... print("\nBasic statistics after cleaning:")
print(train_df.describe())
```

Basic statistics after cleaning:

	date	ad_spend	units \
count	101488	101488.000000	101488.000000
mean	2023-07-09 19:09:44.297651200	84.374182	8.511105
min	2022-04-12 00:00:00	0.000000	0.000000
25%	2023-02-26 00:00:00	0.000000	0.000000
50%	2023-07-16 00:00:00	0.720000	0.000000
75%	2023-12-13 00:00:00	21.642500	3.000000
max	2024-05-31 00:00:00	47934.990000	9004.000000
std	NaN	464.354660	62.680788

	unit_price	missing_name
count	101488.000000	101488.000000
mean	107.256203	0.018032
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	21557.390000	1.000000
std	424.513596	0.133067

In [119... `# CLEANING TEST DATA`

In [120... `test_df.head()`

Out [120...

	ID	date	Item Id	Item Name	ad_spend	anarix_id	unit_
0	2024-07-01_B09KDR64LT	2024-07-01	B09KDR64LT	NapQueen Elizabeth 10" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
1	2024-07-01_B09KDTS4DC	2024-07-01	B09KDTS4DC	NapQueen Elizabeth 8" Gel Memory Foam Mattress...	NaN	NAPQUEEN	
2	2024-07-01_B09KDTHJ6V	2024-07-01	B09KDTHJ6V	NapQueen Elizabeth 12" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
3	2024-07-01_B09KDQ2BWY	2024-07-01	B09KDQ2BWY	NapQueen Elizabeth 12" Gel Memory Foam Mattres...	NaN	NAPQUEEN	
4	2024-07-01_B09KDYY3SB	2024-07-01	B09KDYY3SB	NapQueen Elizabeth 10" Gel Memory Foam Mattres...	101.72	NAPQUEEN	1

```
In [121...] test_df['date'] = pd.to_datetime(test_df['date'])
```

```
In [122...] test_df['ad_spend'] = test_df['ad_spend'].fillna(0)
```

```
In [123...] test_df['Item Name'] = test_df['Item Name'].fillna('Unknown')
```

```
In [124...] test_df['unit_price'] = test_df['unit_price'].clip(lower=0)
test_df['missing_name'] = test_df['Item Name'].apply(lambda x: 1 if x ==
```

```
In [125...] print("Rows in dataset after cleaning:", len(test_df))
print("\nMissing values after cleaning:")
print(test_df.isnull().sum())
print("\nDataset info after cleaning:")
test_df.info()
```

Rows in dataset after cleaning: 2833

Missing values after cleaning:

```
ID          0
date         0
Item Id      0
Item Name    0
ad_spend     0
anarix_id    0
unit_price   0
missing_name 0
dtype: int64
```

Dataset info after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2833 entries, 0 to 2832
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	2833 non-null	object
1	date	2833 non-null	datetime64[ns]
2	Item Id	2833 non-null	object
3	Item Name	2833 non-null	object
4	ad_spend	2833 non-null	float64
5	anarix_id	2833 non-null	object
6	unit_price	2833 non-null	float64
7	missing_name	2833 non-null	int64

```
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
```

```
memory usage: 177.2+ KB
```

```
In [126... print("\nBasic statistics after cleaning:")
           print(test_df.describe())
```

Basic statistics after cleaning:

	date	ad_spend	unit_price	missing_n
count	2833	2833.000000	2833.000000	2833.000000
mean	2024-07-13 16:54:33.420402176	96.997586	100.667754	0.121426
min	2024-07-01 00:00:00	0.000000	0.000000	0.000000
25%	2024-07-07 00:00:00	0.000000	0.000000	0.000000
50%	2024-07-14 00:00:00	0.000000	0.000000	0.000000
75%	2024-07-20 00:00:00	36.030000	0.000000	0.000000
max	2024-07-28 00:00:00	18724.850000	6870.000000	1.000000
std	NaN	565.607189	379.505541	0.326679

STEP 4: EDA (Exploratory Data Analysis)

```
In [127... # Basic info
print("Train dataset shape:", train_df.shape)
print("Test dataset shape:", test_df.shape)

print("\nTrain dataset columns:")
print(train_df.columns)
print("\nTest dataset columns:")
print(test_df.columns)

print("\nTrain dataset info:")
train_df.info()

print("\nTest dataset info:")
test_df.info()
```

Train dataset shape: (101488, 9)

Test dataset shape: (2833, 8)

Train dataset columns:

```
Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id', 'units',
      'unit_price', 'missing_name'],
      dtype='object')
```

Test dataset columns:

```
Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id',
      'unit_price', 'missing_name'],
      dtype='object')
```

Train dataset info:

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 101488 entries, 62967 to 101466

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	ID	101488 non-null	object
1	date	101488 non-null	datetime64[ns]
2	Item Id	101488 non-null	object
3	Item Name	101488 non-null	object
4	ad_spend	101488 non-null	float64
5	anarix_id	101488 non-null	object
6	units	101488 non-null	float64
7	unit_price	101488 non-null	float64
8	missing_name	101488 non-null	int64

dtypes: datetime64[ns](1), float64(3), int64(1), object(4)

memory usage: 7.7+ MB

Test dataset info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 2833 entries, 0 to 2832

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	ID	2833 non-null	object
1	date	2833 non-null	datetime64[ns]
2	Item Id	2833 non-null	object
3	Item Name	2833 non-null	object
4	ad_spend	2833 non-null	float64
5	anarix_id	2833 non-null	object
6	unit_price	2833 non-null	float64
7	missing_name	2833 non-null	int64

dtypes: datetime64[ns](1), float64(2), int64(1), object(4)

memory usage: 177.2+ KB

```
In [128... print(train_df.describe())

# Check for missing values
print(train_df.isnull().sum())
print(test_df.isnull().sum())
```


	date	ad_spend	units \
count	101488	101488.000000	101488.000000
mean	2023-07-09 19:09:44.297651200	84.374182	8.511105
min	2022-04-12 00:00:00	0.000000	0.000000
25%	2023-02-26 00:00:00	0.000000	0.000000
50%	2023-07-16 00:00:00	0.720000	0.000000
75%	2023-12-13 00:00:00	21.642500	3.000000
max	2024-05-31 00:00:00	47934.990000	9004.000000
std	NaN	464.354660	62.680788

	unit_price	missing_name
count	101488.000000	101488.000000
mean	107.256203	0.018032
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	21557.390000	1.000000
std	424.513596	0.133067

```
ID      0
date    0
Item Id  0
Item Name 0
ad_spend 0
anarix_id 0
units    0
unit_price 0
missing_name 0
dtype: int64
ID      0
date    0
Item Id  0
Item Name 0
ad_spend 0
anarix_id 0
unit_price 0
missing_name 0
dtype: int64
```

```
In [129... train_df['date'] = pd.to_datetime(train_df['date'])
test_df['date'] = pd.to_datetime(test_df['date'])

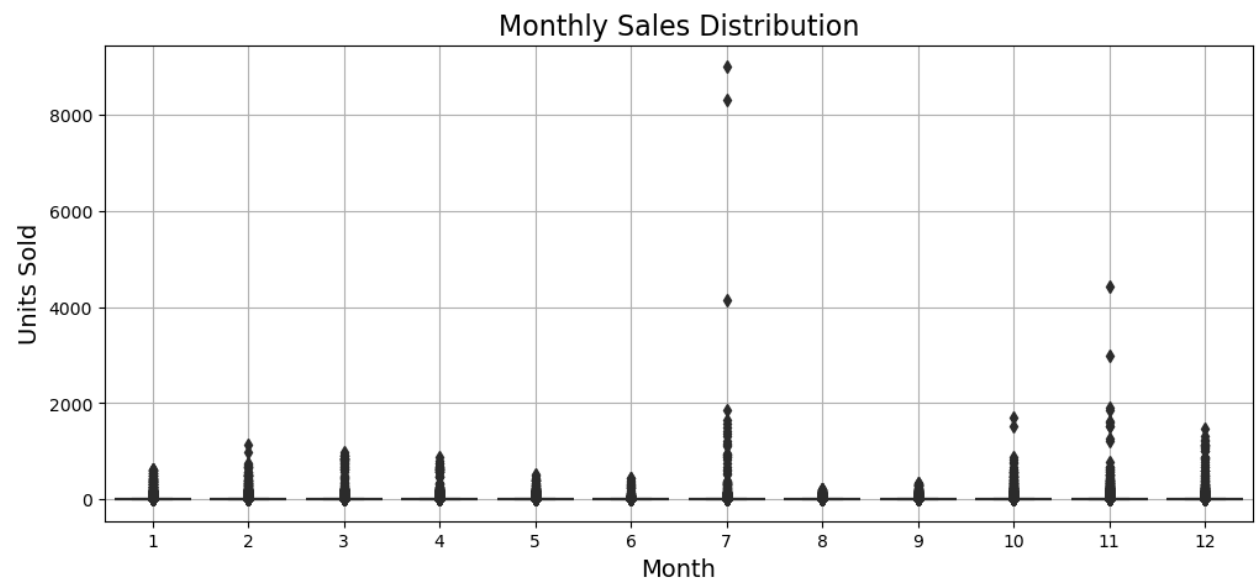
plt.figure(figsize=(15,5))
plt.plot(train_df.groupby('date')['units'].sum(), color='teal', linestyle
plt.title('Total Units Sold Over Time', fontsize=16)
plt.xlabel('Date', fontsize=14)
plt.ylabel('Total Units Sold', fontsize=14)
plt.grid(True)
plt.show()
```

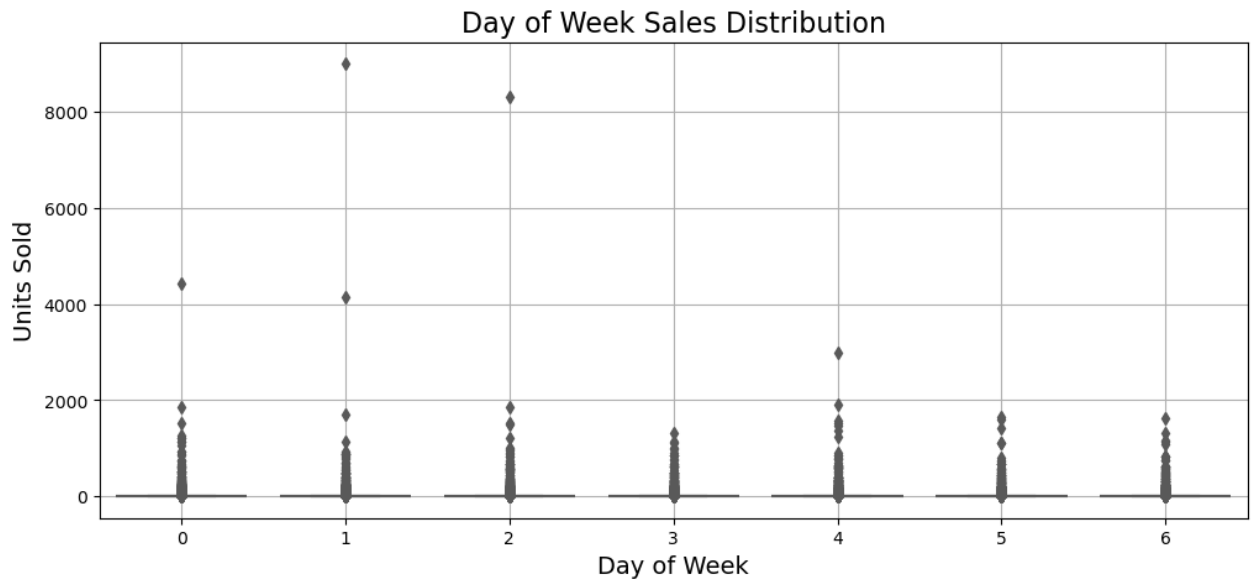


```
In [130... #seasonality
train_df['month'] = train_df['date'].dt.month
train_df['day_of_week'] = train_df['date'].dt.dayofweek


plt.figure(figsize=(12,5))
sns.boxplot(x='month', y='units', data=train_df, palette='Blues')
plt.title('Monthly Sales Distribution', fontsize=16)
plt.xlabel('Month', fontsize=14)
plt.ylabel('Units Sold', fontsize=14)
plt.grid(True)
plt.show()

plt.figure(figsize=(12,5))
sns.boxplot(x='day_of_week', y='units', data=train_df, palette='coolwarm')
plt.title('Day of Week Sales Distribution', fontsize=16)
plt.xlabel('Day of Week', fontsize=14)
plt.ylabel('Units Sold', fontsize=14)
plt.grid(True)
plt.show()
```





```
In [131]: # correlation analysis
corr = train_df[['ad_spend', 'units', 'unit_price']].corr()
sns.heatmap(corr, annot=True)
plt.title('Correlation Heatmap')
plt.show()
```



```
In [132]: #item analysis
#top selling items
top_items = train_df.groupby('Item Id')['units'].sum().nlargest(10)
print("Top 10 selling items:", top_items)
```

```
#items with highest ad spend
top_ad_spend = train_df.groupby('Item Id')['ad_spend'].sum().nlargest(10)
print("\nTop 10 items by ad spend:", top_ad_spend)
```

Top 10 selling items: Item Id

B09KT5HMNY	102149.0
B0BGDWGRQB	99586.0
B09MR3Y296	82697.0
B09MR4B13C	78548.0
B0BGDX2Z3L	33810.0
B09KTMKDKJ	31790.0
B0BRCYQNSW	26531.0
B09MR3XT5G	21815.0
B0B699PLXD	21622.0
B0BGDZLDX2	20691.0

Name: units, dtype: float64

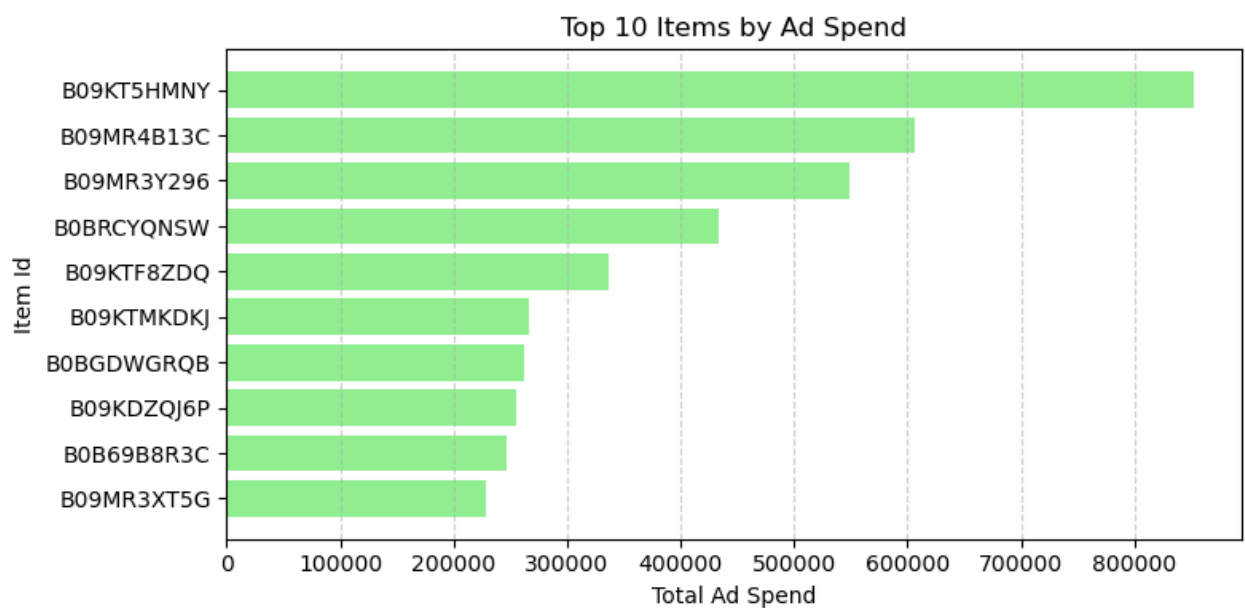
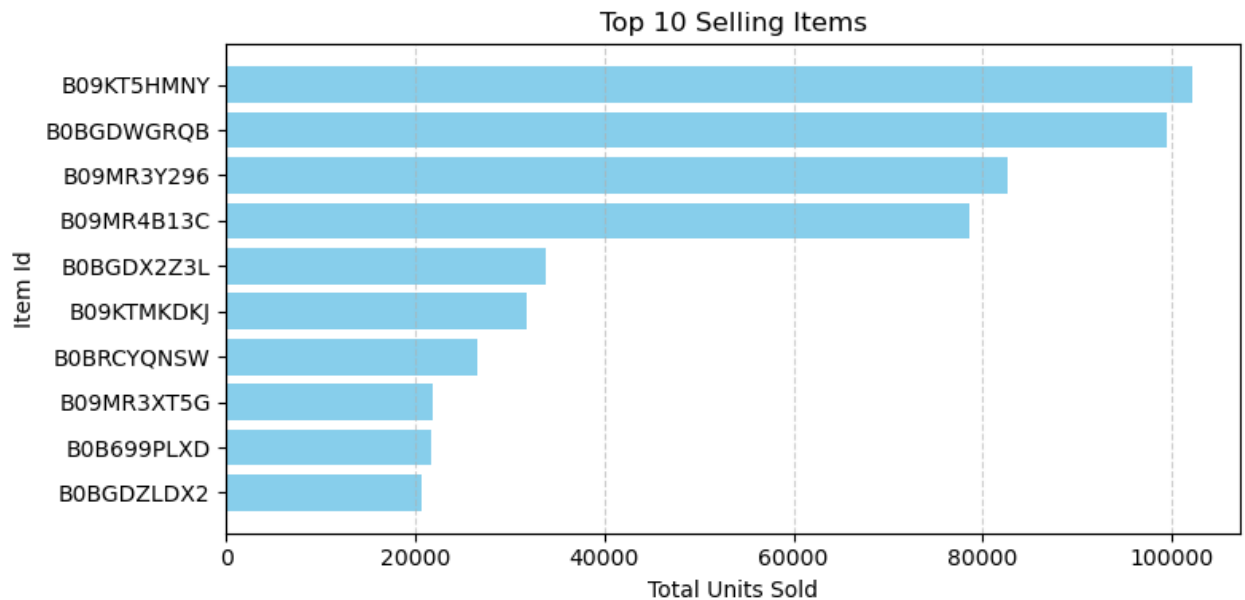
Top 10 items by ad spend: Item Id

B09KT5HMNY	851785.25
B09MR4B13C	607033.59
B09MR3Y296	548484.97
B0BRCYQNSW	434229.75
B09KTF8ZDQ	336946.80
B09KTMKDKJ	266972.17
B0BGDWGRQB	262212.18
B09KDZQJ6P	255192.82
B0B69B8R3C	247132.75
B09MR3XT5G	227750.13

Name: ad_spend, dtype: float64

```
In [133... #horizontal bar graph for top selling items
plt.figure(figsize=(8, 4))
plt.barh(top_items.index, top_items.values, color='skyblue')
plt.xlabel('Total Units Sold')
plt.ylabel('Item Id')
plt.title('Top 10 Selling Items')
plt.gca().invert_yaxis() # Invert y-axis to display top item at the top
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

#horizontal bar graph for items with highest ad spend
plt.figure(figsize=(8, 4))
plt.barh(top_ad_spend.index, top_ad_spend.values, color='lightgreen')
plt.xlabel('Total Ad Spend')
plt.ylabel('Item Id')
plt.title('Top 10 Items by Ad Spend')
plt.gca().invert_yaxis() # Invert y-axis to display top item at the top
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```





```
In [134... #unitprice over time (train and test split)

train_df['date'] = pd.to_datetime(train_df['date'])
test_df['date'] = pd.to_datetime(test_df['date'])

fig, ax = plt.subplots(figsize=(15, 5))
train_df.plot(x='date', y='unit_price', ax=ax, label='Training Set', color='blue')
test_df.plot(x='date', y='unit_price', ax=ax, label='Test Set', color='orange')
ax.axvline(pd.Timestamp('2024-07-01'), color='black', linestyle='--')
ax.set_title('Unit Price Over Time: Train/Test Split')
ax.set_xlabel('Date')
ax.set_ylabel('Unit Price')
ax.legend(['Training Set', 'Test Set', 'Train/Test Split'])

plt.show()
```



STEP 5: FEATURE CREATION

```
In [135... print(train_df.columns)
```

```
Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id', 'units',
      'unit_price', 'missing_name', 'month', 'day_of_week'],
      dtype='object')
```

```
In [136... print(test_df.columns)
```

```
Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id',
      'unit_price', 'missing_name'],
      dtype='object')
```

```
In [137... train_df['date'] = pd.to_datetime(train_df['date'])
test_df['date'] = pd.to_datetime(test_df['date'])
```

```
In [138... def create_features(df):
    if 'date' not in df.columns:
        raise ValueError("DataFrame must contain a 'date' column")

    df['date'] = pd.to_datetime(df['date'])

    df['day_of_week'] = df['date'].dt.dayofweek
    df['month'] = df['date'].dt.month
    df['quarter'] = df['date'].dt.quarter
    df['year'] = df['date'].dt.year
    df['is_weekend'] = df['day_of_week'].isin([5, 6]).astype(int)
    return df

train_df = create_features(train_df)
```

```
In [139... train_df = create_features(train_df)
test_df = create_features(test_df)
```

```
In [140... features = ['day_of_week', 'month', 'quarter', 'year', 'is_weekend']
```

```
In [141... print(train_df.columns)

Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id', 'units',
       'unit_price', 'missing_name', 'month', 'day_of_week', 'quarter', 'year',
       'is_weekend'],
      dtype='object')
```

```
In [142... print(test_df.columns)

Index(['ID', 'date', 'Item Id', 'Item Name', 'ad_spend', 'anarix_id',
       'unit_price', 'missing_name', 'day_of_week', 'month', 'quarter', 'year',
       'is_weekend'],
      dtype='object')
```

```
In [143... column_to_drop = ['ad_spend']
train_df = train_df.drop(columns=column_to_drop, errors='ignore')
test_df = test_df.drop(columns=column_to_drop, errors='ignore')

print(f"Train dataset shape: {train_df.shape}")
print(f"Test dataset shape: {test_df.shape}")
```

Train dataset shape: (101488, 13)

Test dataset shape: (2833, 12)

```
In [144... print("Shape of dataset after feature engineering:", train_df.shape)
print("\nNew features added:")
print(train_df[['month', 'day_of_week', 'quarter', 'year',
                 'is_weekend']].head())
```

Shape of dataset after feature engineering: (101488, 13)

New features added:

	month	day_of_week	quarter	year	is_weekend
62967	9	0	3	2023	0
65359	10	1	4	2023	0
69206	11	3	4	2023	0
767	5	3	2	2022	0
812	5	4	2	2022	0

```
In [145... train_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 101488 entries, 62967 to 101466
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ID                     101488 non-null object  
1   date                   101488 non-null datetime64[ns]
2   Item Id                101488 non-null object  
3   Item Name              101488 non-null object  
4   anarix_id              101488 non-null object  
5   units                  101488 non-null float64  
6   unit_price              101488 non-null float64  
7   missing_name           101488 non-null int64  
8   month                  101488 non-null int32  
9   day_of_week            101488 non-null int32  
10  quarter                 101488 non-null int32  
11  year                    101488 non-null int32  
12  is_weekend              101488 non-null int64  
dtypes: datetime64[ns](1), float64(2), int32(4), int64(2), object(4)
memory usage: 9.3+ MB

```

In [146... train_df.describe()

Out [146...

	date	units	unit_price	missing_name	
count	101488	101488.000000	101488.000000	101488.000000	101488.0
mean	2023-07-09 19:09:44.297651200	8.511105	107.256203	0.018032	6.0
min	2022-04-12 00:00:00	0.000000	0.000000	0.000000	1.0
25%	2023-02-26 00:00:00	0.000000	0.000000	0.000000	3.0
50%	2023-07-16 00:00:00	0.000000	0.000000	0.000000	6.0
75%	2023-12-13 00:00:00	3.000000	0.000000	0.000000	9.0
max	2024-05-31 00:00:00	9004.000000	21557.390000	1.000000	12.0
std	NaN	62.680788	424.513596	0.133067	3.4

In [147... train_df.isnull().sum()

```
Out[147... ID          0
          date        0
          Item Id     0
          Item Name    0
          anarix_id    0
          units        0
          unit_price   0
          missing_name 0
          month        0
          day_of_week  0
          quarter      0
          year         0
          is_weekend   0
          dtype: int64
```

```
In [148... test_df.isnull().sum()
```

```
Out[148... ID          0
          date        0
          Item Id     0
          Item Name    0
          anarix_id    0
          unit_price   0
          missing_name 0
          day_of_week  0
          month        0
          quarter      0
          year         0
          is_weekend   0
          dtype: int64
```

```
In [149... print("Shape of dataset after feature engineering:", train_df.shape)
          print("\nNew features added:")
          print(train_df[['month', 'day_of_week', 'quarter', 'year']].head())
```

Shape of dataset after feature engineering: (101488, 13)

New features added:

	month	day_of_week	quarter	year
62967	9	0	3	2023
65359	10	1	4	2023
69206	11	3	4	2023
767	5	3	2	2022
812	5	4	2	2022

```
In [150... print("Shape of dataset after feature engineering:", test_df.shape)
          print("\nNew features added:")
          print(test_df[['month', 'day_of_week', 'quarter', 'year',
                        'is_weekend']].head())
```

Shape of dataset after feature engineering: (2833, 12)

New features added:

	month	day_of_week	quarter	year	is_weekend
0	7	0	3	2024	0
1	7	0	3	2024	0
2	7	0	3	2024	0
3	7	0	3	2024	0
4	7	0	3	2024	0

STEP 6: MODEL SELECTION

```
In [151...] train_df['date'] = pd.to_datetime(train_df['date'])
test_df['date'] = pd.to_datetime(test_df['date'])
```

b. XGBoost

```
In [152...] print(train_df.columns)

Index(['ID', 'date', 'Item Id', 'Item Name', 'anarix_id', 'units',
       'unit_price', 'missing_name', 'month', 'day_of_week', 'quarter', 'year',
       'is_weekend'],
      dtype='object')
```

```
In [153...] print(test_df.columns)

Index(['ID', 'date', 'Item Id', 'Item Name', 'anarix_id', 'unit_price',
       'missing_name', 'day_of_week', 'month', 'quarter', 'year',
       'is_weekend'],
      dtype='object')
```

```
In [154...] # Separate features and target variable from training data
X_train = train_df[features]
y_train = train_df['units']
```

```
In [155...] X_test = test_df[features]
```

```
In [156...] model = xgb.XGBRegressor(
    subsample=0.7,
    n_estimators=1000,
    min_child_weight=3,
    max_depth=3,
    learning_rate=0.01
)
```

```
In [157...] model.fit(X_train, y_train)
```

Out [157...

▼ XGBRegressor

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_r
ounds=None,
              enable_categorical=False, eval_metric=None, feature_
types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.01, ma
x_bin=None,
```

In [158...

```
model.fit(
    X_train, y_train,
    eval_set=[(X_train, y_train)],
    verbose=100
)
```

```
[0]    validation_0-rmse:62.67904
[100]   validation_0-rmse:62.61096
[200]   validation_0-rmse:62.58460
[300]   validation_0-rmse:62.57098
[400]   validation_0-rmse:62.56220
[500]   validation_0-rmse:62.55689
[600]   validation_0-rmse:62.55265
[700]   validation_0-rmse:62.54844
[800]   validation_0-rmse:62.54572
[900]   validation_0-rmse:62.54352
[999]   validation_0-rmse:62.54133
```

Out [158...

▼ XGBRegressor

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_r
ounds=None,
              enable_categorical=False, eval_metric=None, feature_
types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.01, ma
x_bin=None,
```

In [159...

```
test_df['units_pred'] = model.predict(X_test)
```

In [160...


```
submission = test_df[['ID']]
submission['Target'] = test_df['units_pred']
```

```
In [161... submission.to_csv('submission.csv', index=False)
print("Submission file created successfully!")
```

Submission file created successfully!

```
In [162... feature_importance = model.feature_importances_
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': feature_importance
})
importance_df = importance_df.sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'], importance_df['Importance'])
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Feature Importance')
plt.show()
```



```
In [163... from sklearn.metrics import mean_squared_error
train_predictions = model.predict(X_train)

mse = mean_squared_error(y_train, train_predictions)

print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 3911.418409604966

PERFORMING HYPER PARAMETER TUNING


```
In [164... import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error
from xgboost import XGBRegressor
```

```
In [165... features = ['day_of_week', 'month', 'quarter', 'year', 'is_weekend']
target = 'units'
```

```
In [166... X_train = train_df[features]
y_train = train_df[target]
```

```
In [167... X_train_split, X_val, y_train_split, y_val = train_test_split(X_train, y_
```

```
In [168... xgb = XGBRegressor()
```

```
In [169... # Hyperparameter tuning
param_grid = {
    'n_estimators': [100, 200],
    'learning_rate': [0.01, 0.1],
    'max_depth': [3, 5],
    'min_child_weight': [1, 3],
    'subsample': [0.7, 1.0],
    'colsample_bytree': [0.7, 1.0]
}
```

```
In [170... grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, scoring=
grid_search.fit(X_train_split, y_train_split)
```

```
Fitting 3 folds for each of 64 candidates, totalling 192 fits
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=200, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
weight=1, n_estimators=200, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=0.7, learning_rate=0.01, max_depth=3, min_child_
```

Page 26 of 33

Page 27 of 33

Page 28 of 33

Page 29 of 33

[illegible]

Page 31 of 33

```

eight=1, n_estimators=100, subsample=1.0; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=200, subsample=0.7; total time= 0.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=200, subsample=1.0; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=1, n_estimators=200, subsample=1.0; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=100, subsample=0.7; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=100, subsample=1.0; total time= 0.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=0.7; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=1.0; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=1.0; total time= 0.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, min_child_w
eight=3, n_estimators=200, subsample=1.0; total time= 0.2s

```

Out [170]...

```

▶ GridSearchCV
▶ estimator: XGBRegressor
  ▶ XGBRegressor

```

In [171]...

```

# Best parameters and model
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

```

In [172]...

```

# Evaluate on validation set
val_predictions = best_model.predict(X_val)
mse = mean_squared_error(y_val, val_predictions)

```



```
print(f"Best parameters: {best_params}")
print(f"Validation MSE: {mse}")
```

Best parameters: {'colsample_bytree': 0.7, 'learning_rate': 0.01, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 1.0}
Validation MSE: 1443.7185930621974

```
In [173...] best_model.fit(X_train, y_train)
```

```
Out[173...] XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.7, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.01, max_bin=None,
```

```
In [174...] #test data predictions
X_test = test_df[features]
test_predictions = best_model.predict(X_test)
```

```
In [176...] submission = pd.DataFrame({'ID': test_df['ID'], 'Target': test_predictions})
submission.to_csv('submission2.csv', index=False)
```

```
In [ ]:
```